

Performance and Languages

Troels Henriksen

22nd of September, 2021

What is performance?

What is performance?

- How **fast** you solve a problem.
- How **few resources** you use to solve a problem.

What is performance?

- How **fast** you solve a problem.
- How **few resources** you use to solve a problem.

Important: the program must still be correct—it's easy to make a very fast program if it doesn't have to produce the right result.

What is performance?

- How **fast** you solve a problem.
- How **few resources** you use to solve a problem.

Important: the program must still be correct—it's easy to make a very fast program if it doesn't have to produce the right result.

There's more to performance than asymptotic complexity!

Constant factors matter

- Easily see 10:1 performance range depending on how code is written
- Must optimize at multiple levels:
 - ▶ algorithm, data representations, procedures, and loops
- Must understand system to optimise performance:
 - ▶ How programs are compiled and executed
 - ▶ How modern processors + memory systems operate
 - ▶ How to measure program performance and identify bottlenecks
 - ▶ How to improve performance without destroying code modularity and generality

Constant factors matter

- Easily see 10:1 performance range depending on how code is written
- Must optimize at multiple levels:
 - ▶ algorithm, data representations, procedures, and loops
- Must understand system to optimise performance:
 - ▶ How programs are compiled and executed
 - ▶ How modern processors + memory systems operate
 - ▶ How to measure program performance and identify bottlenecks
 - ▶ How to improve performance without destroying code modularity and generality

Let's look at why C is often considered a “fast language”, and what that means.

Roughly two kinds of languages

Compiled languages are transformed to machine code before execution (e.g. C)

Interpreted languages are run directly by a software *interpreter* (e.g. Python)

Roughly two kinds of languages

Compiled languages are transformed to machine code before execution (e.g. C)

Interpreted languages are run directly by a software *interpreter* (e.g. Python)

Pedantic disclaimer

Compilation/interpretation is strictly a property of *implementations*, not *languages*.

- You could have a C interpreter or Python compiler
- But most (not all!) languages are built with a specific implementation technique in mind
- A few languages (Lisp, JavaScript) have lots of *very* different implementations

Roughly two kinds of languages

Compiled languages are transformed to machine code before execution (e.g. C)

Interpreted languages are run directly by a software *interpreter* (e.g. Python)

Pedantic disclaimer

Compilation/interpretation is strictly a property of *implementations*, not *languages*.

- You could have a C interpreter or Python compiler
- But most (not all!) languages are built with a specific implementation technique in mind
- A few languages (Lisp, JavaScript) have lots of *very* different implementations

We teach you the big picture—the details are always more complicated in practice!

Tradeoffs

- Compiled languages

- + Almost always faster
 - Require compilation after every change
 - Usually cannot run program fragments in isolation
 - Tend to have more restrictions (e.g. static typing)
 - Much more difficult to implement

Tradeoffs

- Compiled languages

- + Almost always faster
 - Require compilation after every change
 - Usually cannot run program fragments in isolation
 - Tend to have more restrictions (e.g. static typing)
 - Much more difficult to implement

- Interpreted languages

- Usually slow
- + Can run immediately
- + Can easily run fragments (e.g. single functions) in isolation
- + Much easier to implement

Tradeoffs

- Compiled languages
 - + Almost always faster
 - Require compilation after every change
 - Usually cannot run program fragments in isolation
 - Tend to have more restrictions (e.g. static typing)
 - Much more difficult to implement
- Interpreted languages
 - Usually slow
 - + Can run immediately
 - + Can easily run fragments (e.g. single functions) in isolation
 - + Much easier to implement

Let us look at the scale of the overhead.

The Collatz conjecture

$$f(n) = \left\{ \begin{array}{ll} \frac{n}{2} & \text{if } n \text{ is even} \\ 3n + 1 & \text{if } n \text{ is odd} \end{array} \right\}$$

- **Conjecture:** if we apply this function to some number greater than 1, we will eventually reach 1
- To disprove this conjecture, we only need *a single counter-example* that goes into a cycle instead
- People write programs to investigate the behaviour of this sequence

Listing 1: collatz.py

```
import sys

def collatz(n):
    i = 0
    while n != 1:
        if n % 2 == 0:
            n = n // 2
        else:
            n = 3 * n + 1
        i = i + 1
    return i

k = int(sys.argv[1])
for n in range(1, k):
    print(n, collatz(n))
```

Listing 2: collatz.c

```
#include <stdio.h>
#include <stdlib.h>
#include <assert.h>

int collatz(int n) {
    int i = 0;
    while (n != 1) {
        if (n % 2 == 0) { n = n / 2; }
        else { n = 3 * n + 1; }
        i++;
    }
    return i;
}

int main(int argc, char** argv) {
    assert(argc == 2);
    int k = atoi(argv[1]);
    for (int n = 1; n < k; n++) {
        printf("%d_ %d\n", n, collatz(n));
    }
}
```

Benchmarking collatz.py

```
$ time python3 ./collatz.py 100000 >/dev/null
```

```
real    0m1.368s
```

```
user    0m1.361s
```

```
sys     0m0.007s
```


Benchmarking collatz.py

```
$ time python3 ./collatz.py 100000 >/dev/null
```

```
real    0m1.368s
```

```
user    0m1.361s
```

```
sys     0m0.007s
```

```
$ gcc collatz.c -o collatz
```

Benchmarking collatz.py

```
$ time python3 ./collatz.py 100000 >/dev/null
```

```
real    0m1.368s
```

```
user    0m1.361s
```

```
sys     0m0.007s
```

```
$ gcc collatz.c -o collatz
```

```
$ time ./collatz 100000 >/dev/null
```

```
real    0m0.032s
```

```
user    0m0.030s
```

```
sys     0m0.002s
```

Benchmarking collatz.py

```
$ time python3 ./collatz.py 100000 >/dev/null
```

```
real    0m1.368s
```

```
user    0m1.361s
```

```
sys     0m0.007s
```

```
$ gcc collatz.c -o collatz
```

```
$ time ./collatz 100000 >/dev/null
```

```
real    0m0.032s
```

```
user    0m0.030s
```

```
sys     0m0.002s
```

Speedup: $\frac{1.368}{0.032} = 42.75$

Why are interpreted languages slow?

To understand this, we must look at **how processors execute code**.

Register file
rax
rcx
rdx
rbx
...

- The CPU has a fixed set of **registers**
- Registers act as variables for **machine code instructions**
- Some registers are **general purpose**, others are **special purpose**
 - ▶ Program counter holds address of current instruction
 - ▶ Condition register holds results of comparisons
 - ▶ Stack pointer holds address of top of stack
- Exact registers and instructions depend on the **architecture**.
 - ▶ We teach you x86-64 (sometimes called AMD64)
 - ▶ Others are ARM, MIPS, SPARC, POWER...
 - ▶ On x86-64, most registers hold 64-bit quantities.

Compiling C to assembly

```
$ gcc mul.c -S
```

Listing 3: mul.c

```
long mul(long x, long y) {  
    return x * y;  
}
```

Listing 4: mul.s (extract)

```
mul:  
.LFB0:  
    .cfi_startproc  
    movq    %rdi, %rax  
    imulq   %rsi, %rax  
    ret  
    .cfi_endproc
```

Calling convention:

- The two arguments stored in registers `rdi` and `rsi`.
- Return value stored in register `rax`.

A modern processor can execute billions of instructions per second.

Compiling Python to...?

The Python interpreter actually does *compile* Python to “opcodes”.

```
$ python
```

```
>>> def mul(x,y): return x * y
```

```
>>> import dis
```

```
>>> dis.dis(mul)
```

1	0	LOAD_FAST	0 (x)
	2	LOAD_FAST	1 (y)
	4	BINARY_MULTIPLY	
	6	RETURN_VALUE	

This is “assembly code” for a **virtual machine** that is interpreted by a C program.

Compilation vs interpretation

All programs are interpreted when run!

- Machine code programs (e.g. generated by a C compiler) are run by **an interpreter built in hardware**.
- Python programs are run by **an interpreter written in software**, which must *itself* be interpreted by something else.
- **Each level of interpretation adds overhead**, but usually also *convenience*: productivity and correctness!

Combining interpretation and compilation

- Interpreted languages can be fast when
 - ▶ Most of the run-time is spent waiting data from files or network
 - ▶ They mostly call functions written in faster compiled languages
- **Best of both worlds:** flexibility of interpretation, and speed of C

Different ways to compile

Different ways to compile

To executable program collatz

```
$ gcc collatz.c -o collatz
```

- Can be run directly

Different ways to compile

To executable program `collatz`

```
$ gcc collatz.c -o collatz
```

- Can be run directly

To object file `collatz.o`

```
$ gcc collatz.c -c -o collatz.o
```

- Can be *linked* with other object files
- Can be processed further

Different ways to compile

To executable program `collatz`

```
$ gcc collatz.c -o collatz
```

- Can be run directly

To object file `collatz.o`

```
$ gcc collatz.c -c -o collatz.o
```

- Can be *linked* with other object files
- Can be processed further

To shared object file `libcollatz.so`

```
$ gcc collatz.c -fPIC -shared -o libcollatz.so
```

- Can be linked *at run-time* by a running program
- How compiled programs support dynamic “plug-ins”

All output files contain fully compiled machine code.

Calling C from Python

Compiling C program to shared library

```
$ gcc collatz.c -fPIC -shared -o libcollatz.so
```

Listing 5: collatz-ffi.py

```
import ctypes
import sys

c_lib = ctypes.CDLL('./libcollatz.so')

k = int(sys.argv[1])
for n in range(1, k):
    print(n, c_lib.collatz(n))
```

```
$ time python3 ./collatz-ffi.py 100000 >/dev/null
```

```
real    0m0.165s  
user    0m0.163s  
sys     0m0.003s
```

$$\text{Speedup: } \frac{1.368}{0.165} = 8.2$$

```
$ time python3 ./collatz-ffi.py 100000 >/dev/null
```

```
real    0m0.165s
user    0m0.163s
sys     0m0.003s
```

Speedup: $\frac{1.368}{0.165} = 8.2$

- Slower than pure C by about $5\times$
- Faster if we made fewer “foreign” calls, but each took more time
- Ideal case is single foreign function call that operates on many values
- **This is exactly how NumPy works!**

NumPy performance

```
def f_python(v):  
    for i in range(len(v)):  
        v[i] = v[i]*2 + 3
```

```
def f_numpy(v):  
    return v * 2 + 3
```

Size of v	f_python	f_numpy	Difference
1	0.01ms	0.01ms	0.9×
10	0.01ms	0.01ms	1.4×
100	0.1ms	0.01ms	13.3×
1000	0.98ms	0.01ms	95.3×
10000	9.96ms	0.05ms	190.7×
100000	98.59ms	0.41ms	240.7×

Conclusions

- Compiled languages tend to be fast, but less flexible
- Interpreted languages tend to be slower, but more flexible
- **Best of both worlds:** write computational primitives in fast languages, call them from slow languages
 - ▶ NumPy works like this